

Distributed Systems

I5-440 / I5-640

Fall 2015

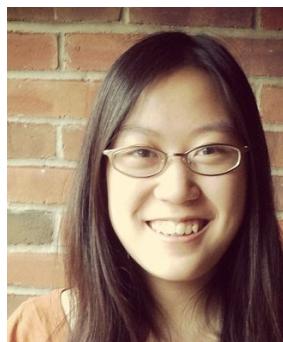
Welcome! Course Staff



Instructors



Aaron Friedlander



Esther Wang



Marisa Chang



Joshua Gluck

9 TA's



Varun Saravgi



Arjun Puri



Chao Xin



Adhish Ramkumar



Xiaoxiang Wu

Course Logistics

- Course Policies
 - Class web page: <http://www.cs.cmu.edu/~srini/15-440/>
 - Piazza: <https://piazza.com/cmu/fall2015/1544015640>
 - Obligatory discussion of {late days, cheating, etc.}
- Waitlist!
- No recitations this year, extra TA hours if needed (??)
- Office hours / TAs are on class web page (keep checking)
- Go work through the Tour of Go!
 - <https://tour.golang.org/welcome/1>

Waitlist

- Waitlist of unprecedented size. Keep coming to class, because we don't really know yet how it will work out.
- Registered: 79 (15-440) + 58 (15-640)
- Waitlisted: 23 (15-440) + 192 (!!!!!)
- The bad news: Not everyone will get in. We are *by law* limited to physical classroom size. This is not subject to negotiation.
- The plea: Not serious about the class? PLEASE DROP SOON.
- The strategy:
 - Attend class!
 - If class is on immediate graduation path, *have academic advisor email us*.

Processing WL/Enroll

- You will not be able to take the class if:
 - If not taken 213/513 at CMU *before*
 - If you are an UGRAD and lower than a “C” in 213
 - If you are a Grad and lower than a “B-” in 213/513
- Priority order
 - Required: CS Ugrad, MSCS, MSDC, MITS
 - Email from faculty advisor
 - ... then WL rank

Course Goals

- Systems requirement:
 - Learn something about distributed systems in particular;
 - Learn general systems principles (modularity, layering, naming, security, ...)
 - Practice implementing real, larger systems; in teams; must run in nasty environment;
- One consequence: Must pass homeworks, exams, and projects independently as well as in total.
 - Note, if you fail either you will not pass the class

Course Format

- ~26 lectures
- Office hours: Practical issues for implementing projects; general questions and discussion
- 4 projects; 2 solo (p0, p1), 2 person team (p2,p3)
 - P1: Distributed (internet-wide) bitcoin miner
 - P2: Building Tribbler (or something)
 - P3: Project with distributed systems concepts like replication or distributed commit/consensus

Book

- Link to Amazon purchase (new, used, rent) from syllabus page
- Several useful references on web page
- We'll be compiling notes (and these slides) for your use over the course of the semester; based on, but not identical to, prior 15-440 instance

About Projects

- Systems programming somewhat different from what you've done before
 - Low-level (C / GO)
 - Often designed to run indefinitely (error handling must be rock solid)
 - Must be secure - horrible environment
 - Concurrency
 - Interfaces specified by documented protocols
- Office Hours & “System Hacker’s View of Software Engineering”
 - Practical techniques designed to save you time & pain
- WARNING: Many students dropped during project I => started too late!

Collaboration

- Working together important
 - Discuss course material
 - Work on problem debugging
- Parts *must* be your own work
 - Homeworks, midterm, final, solo proj
- Team projects: both students should understand entire project
- What we hate to say: we run cheat checkers...
- Please *do not* put code on *public* repositories
- Partner problems: *address early.*

Late Work

- 10% penalty per day
- Cannot be more than 2 days late
 - (no exceptions after 48 hours of due date/time)
- Usual exceptions: documented medical, emergency, etc.
 - *Talk to us early if there's a problem!*
- Regrade requests in writing to course admin

Why take this course?

- Huge amounts of computing are now distributed...
 - A few years ago, Intel threw its hands up in the air: couldn't increase GHz much more without CPU temperatures reaching solar levels
 - But we can still stuff more transistors (Moore's Law)
 - Result: Multi-core and GPUs.
 - Result 2: Your computer has become a parallel/distributed system. In a decade, it may have 128 cores.
- Oh, yeah, and that whole Internet thing...
 - my phone syncs its calendar with google, which i can get on my desktop with a web browser, ...
 - (That phone has the computing power of a desktop from 10 years ago and communicates wirelessly at a rate 5x faster than the average american home could in 1999.)
 - Stunningly impressive capabilities now seem mundane. But *lots* of great stuff going on under the hood...
 - Most things are distributed, and more each day

If you find yourself ...

- In hollywood....
 - ... rendering videos on clusters of 10s of 1000s of nodes?
 - Or getting terabytes of digital footage from on-location to post-processing?
- On wall street...
 - tanking our economy with powerful simulations running on large clusters of machines
 - For 11 years, the NYSE ran software from cornell systems folks to update trade data
- In biochem...
 - using protein folding models that require supercomputers to run
- In gaming...
 - Writing really bad distributed systems to enable MMOs to crash on a regular basis
- not to mention the obvious places

What Is A Distributed System?

“A collection of independent computers that appears to its users as a single coherent system.”

- Features:
 - No shared memory – message-based communication
 - Each runs its own local OS
 - Heterogeneity
- Ideal: to present a single-system image:
 - The distributed system “looks like” a single computer rather than a collection of separate computers.

Distributed System Characteristics

- To present a single-system image:
 - Hide internal organization, communication details
 - Provide uniform interface
- Easily expandable
 - Adding new computers is hidden from users
- Continuous availability
 - Failures in one component can be covered by other components
- Supported by middleware

Definition of a Distributed System

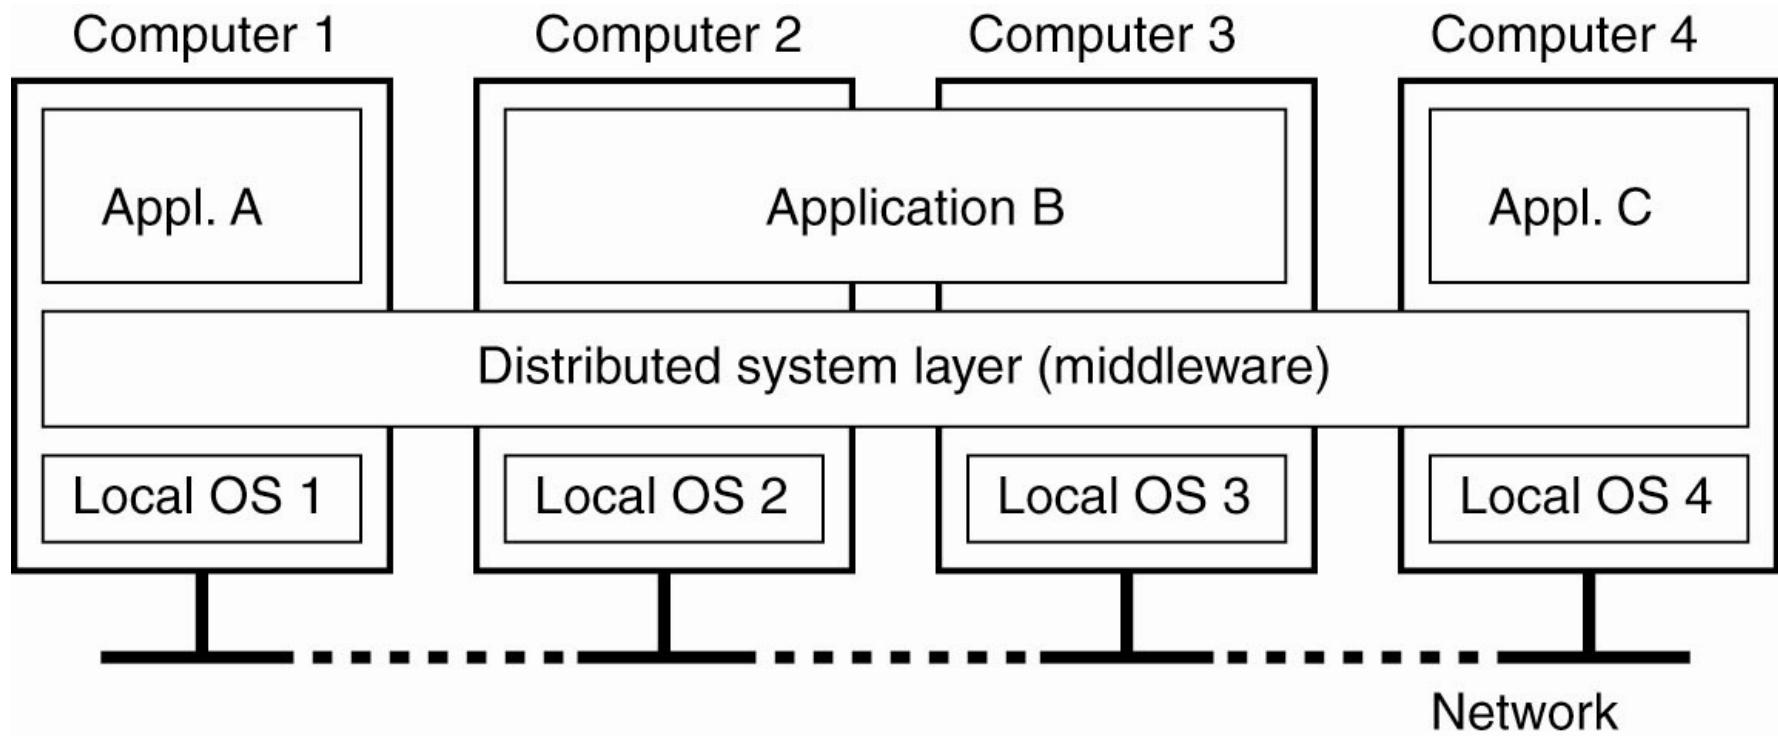


Figure 1-1. A distributed system organized as middleware. The middleware layer runs on all machines, and offers a uniform interface to the system

Goal 1 – Resource Availability

- Support user access to remote resources (printers, data files, web pages, CPU cycles) and the fair sharing of the resources
- Economics of sharing expensive resources
- Performance enhancement – due to multiple processors; also due to ease of collaboration and info exchange – access to remote services
- Resource sharing introduces security problems.

Goal 2 – Distribution Transparency

- Software hides some of the details of the distribution of system resources.
 - Makes the system more user friendly.
- A distributed system that appears to its users & applications to be a single computer system is said to be *transparent*.
 - Users & apps should be able to access remote resources in the same way they access local resources.
- Transparency has several dimensions.

Types of Transparency

Transparency	Description
Access	Hide differences in data representation & resource access (enables interoperability)
Location	Hide location of resource (can use resource without knowing its location)
Migration	Hide possibility that a system may change location of resource (no effect on access)
Replication	Hide the possibility that multiple copies of the resource exist (for reliability and/or availability)
Concurrency	Hide the possibility that the resource may be shared concurrently
Failure	Hide failure and recovery of the resource. How does one differentiate betw. slow and failed?
Relocation	Hide that resource may be moved <u>during use</u>

Figure 1-2. Different forms of transparency in a distributed system (ISO, 1995)

Transparency to Handle Failures?

The Joys of Real Hardware

Typical first year for a new cluster:

- ~0.5 **overheating** (power down most machines in <5 mins, ~1-2 days to recover)
 - ~1 **PDU failure** (~500-1000 machines suddenly disappear, ~6 hours to come back)
 - ~1 **rack-move** (plenty of warning, ~500-1000 machines powered down, ~6 hours)
 - ~1 **network rewiring** (rolling ~5% of machines down over 2-day span)
 - ~20 **rack failures** (40-80 machines instantly disappear, 1-6 hours to get back)
 - ~5 **racks go wonky** (40-80 machines see 50% packetloss)
 - ~8 **network maintenances** (4 might cause ~30-minute random connectivity losses)
 - ~12 **router reloads** (takes out DNS and external vips for a couple minutes)
 - ~3 **router failures** (have to immediately pull traffic for an hour)
 - ~dozens of minor **30-second blips for dns**
 - ~1000 **individual machine failures**
 - ~thousands of **hard drive failures**
- slow disks, bad memory, misconfigured machines, flaky machines, etc.

slide from Jeff Dean, Google

Goal 2: Degrees of Transparency

- Trade-off: transparency versus other factors
- Reduced performance: multiple attempts to contact a remote server can slow down the system – should you report failure and let user cancel request?
- Convenience: direct the print request to my local printer, not one on the next floor
- Too much emphasis on transparency may prevent the user from understanding system behavior.

Goal 3 - Openness

- An **open distributed system** “...offers services according to standard rules that describe the syntax and semantics of those services.” In other words, the interfaces to the system are clearly specified and freely available.
 - Compare to network protocols, *Not* proprietary
- **Interface Definition/Description Languages (IDL):** used to describe the interfaces between software components, usually in a distributed system
 - Definitions are language & machine independent
 - Support communication between systems using different OS/programming languages; e.g. a C++ program running on Windows communicates with a Java program running on UNIX
 - Communication is usually RPC-based.

Examples of IDLs

Goal 3-Openness

- IDL: Interface Description Language
 - The original
 - WSDL: Web Services Description Language
 - Provides machine-readable descriptions of the services
 - OMG IDL: used for RPC in CORBA
 - OMG – Object Management Group
 - ...

Open Systems Support ...

- **Interoperability:** the ability of two different systems or applications to work together
 - A process that needs a service should be able to talk to any process that provides the service.
 - Multiple implementations of the same service may be provided, as long as the interface is maintained
- **Portability:** an application designed to run on one distributed system can run on another system which implements the same interface.
- **Extensibility:** Easy to add new components, features

Goal 4 - Scalability

- Dimensions that may scale:
 - With respect to size
 - With respect to geographical distribution
 - With respect to the number of administrative organizations spanned
- A scalable system still performs well as it scales up along any of the three dimensions.

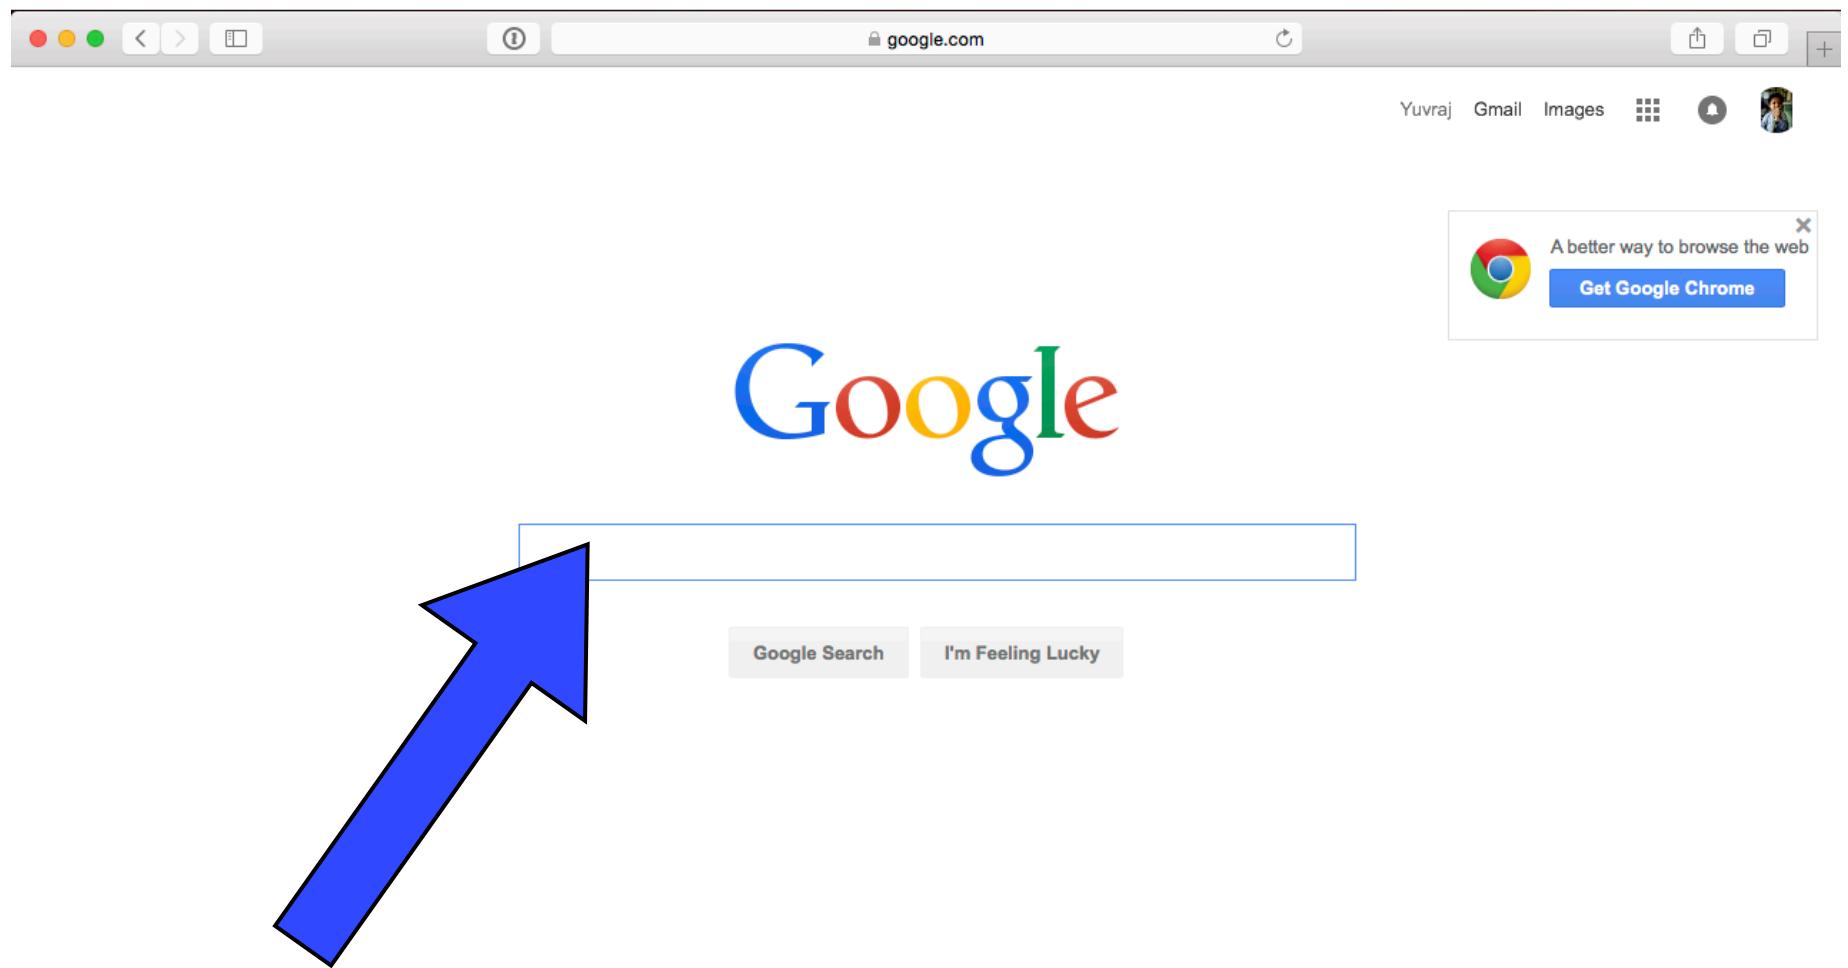
Summary

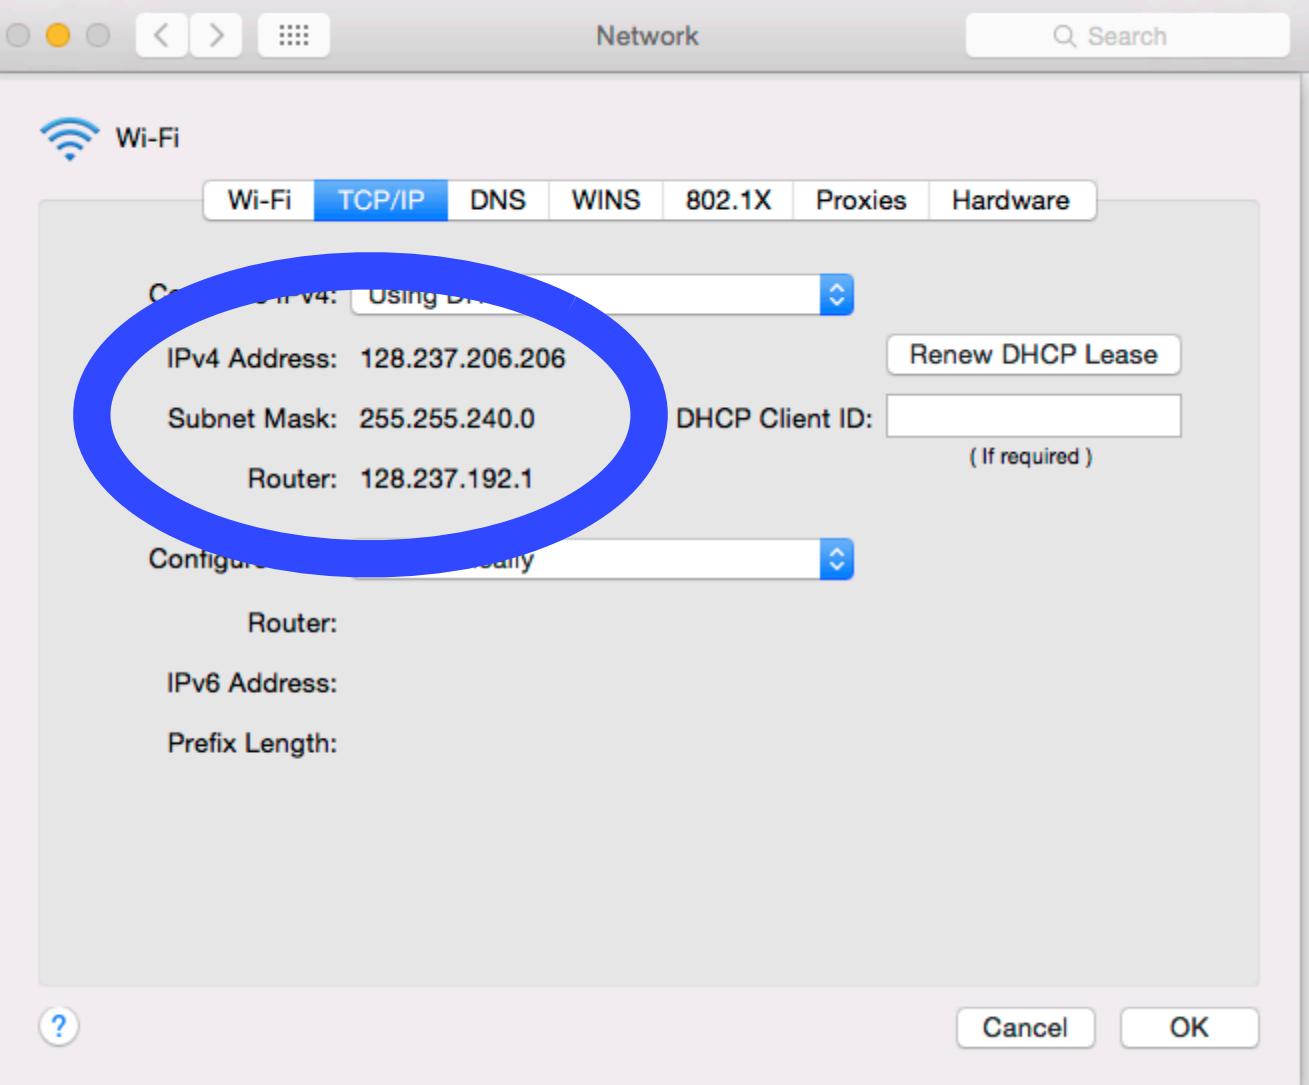
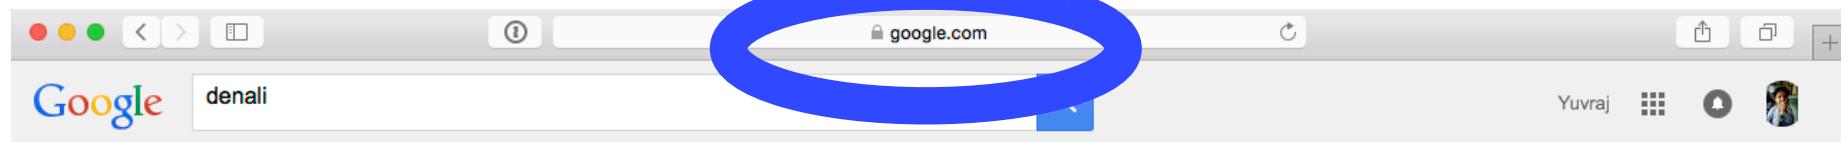
Goals for Distribution

- Resource accessibility
 - For sharing and enhanced performance
- Distribution transparency
 - For easier use
- Openness
 - To support interoperability, portability, extensibility
- Scalability
 - With respect to size (number of users), geographic distribution, administrative domains

Enough advertising

- Let's look at one real distributed system
- That's drastically more complex than it might seem from the web browser...

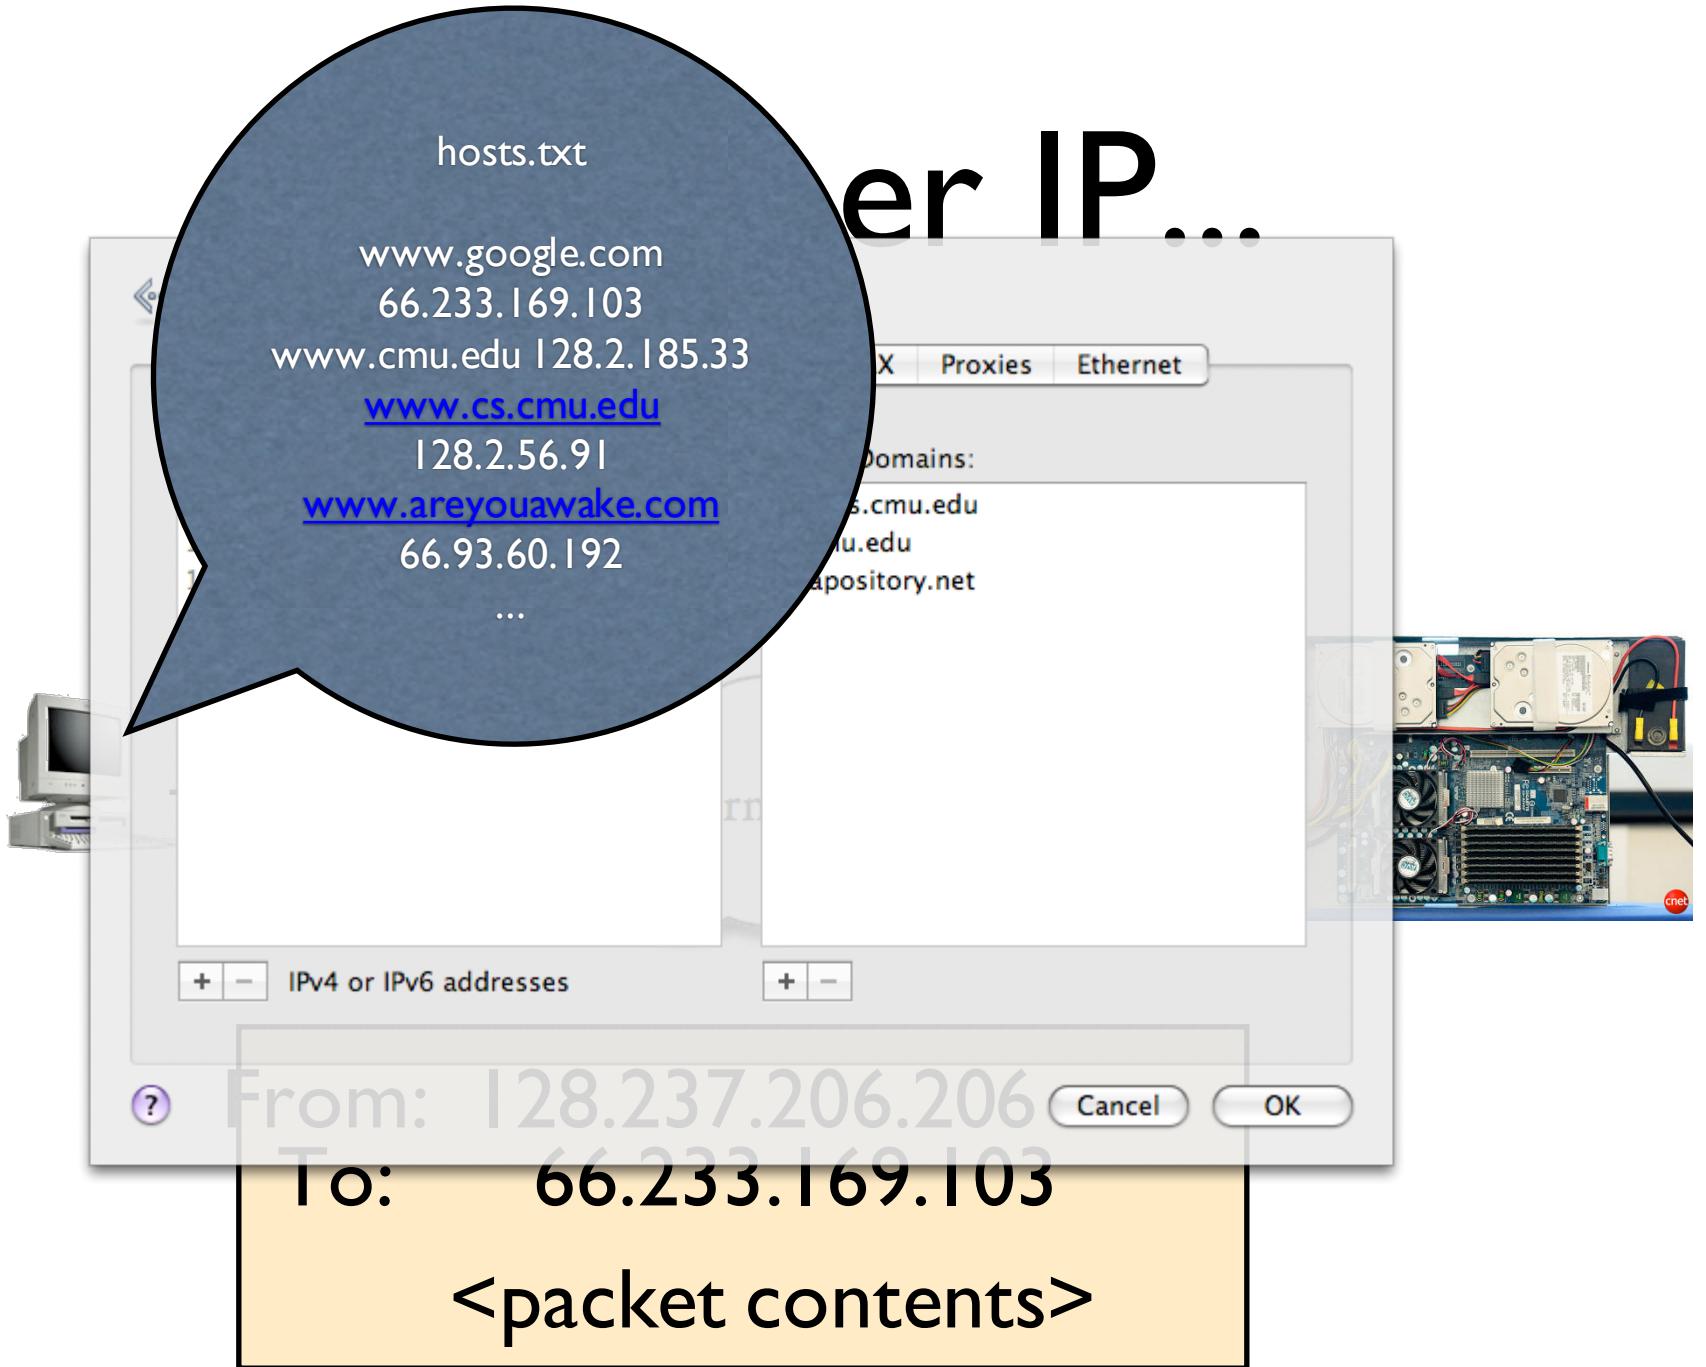




A screenshot of a Mac OS X Network preferences window. The window title is "Network". The "Wi-Fi" tab is selected. The "TCP/IP" tab is highlighted with a blue oval. Other tabs include DNS, WINS, 802.1X, Proxies, and Hardware. Under the TCP/IP tab, the configuration is set to "Using DHCP". The displayed network information includes:

- IPv4 Address: 128.237.206.206
- Subnet Mask: 255.255.240.0
- Router: 128.237.192.1

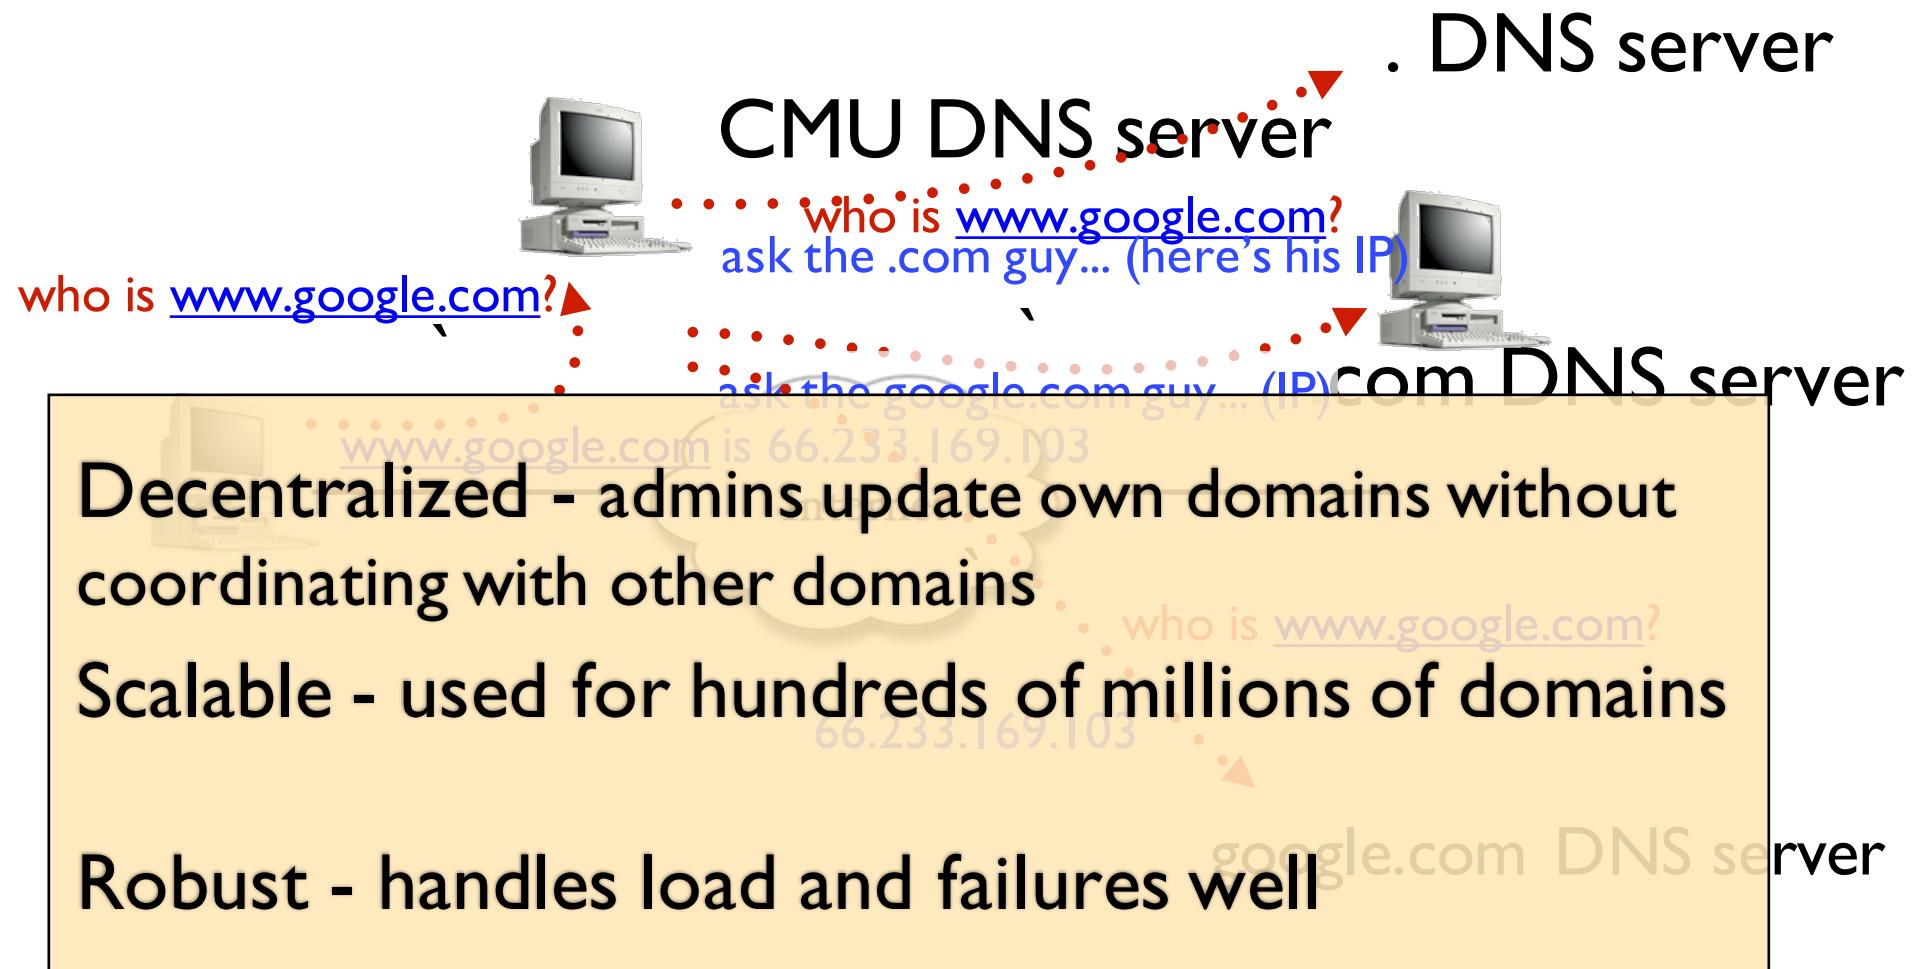
Buttons in the window include "Renew DHCP Lease" and "DHCP Client ID: (If required)". At the bottom are "Cancel" and "OK" buttons. A question mark icon is also present.



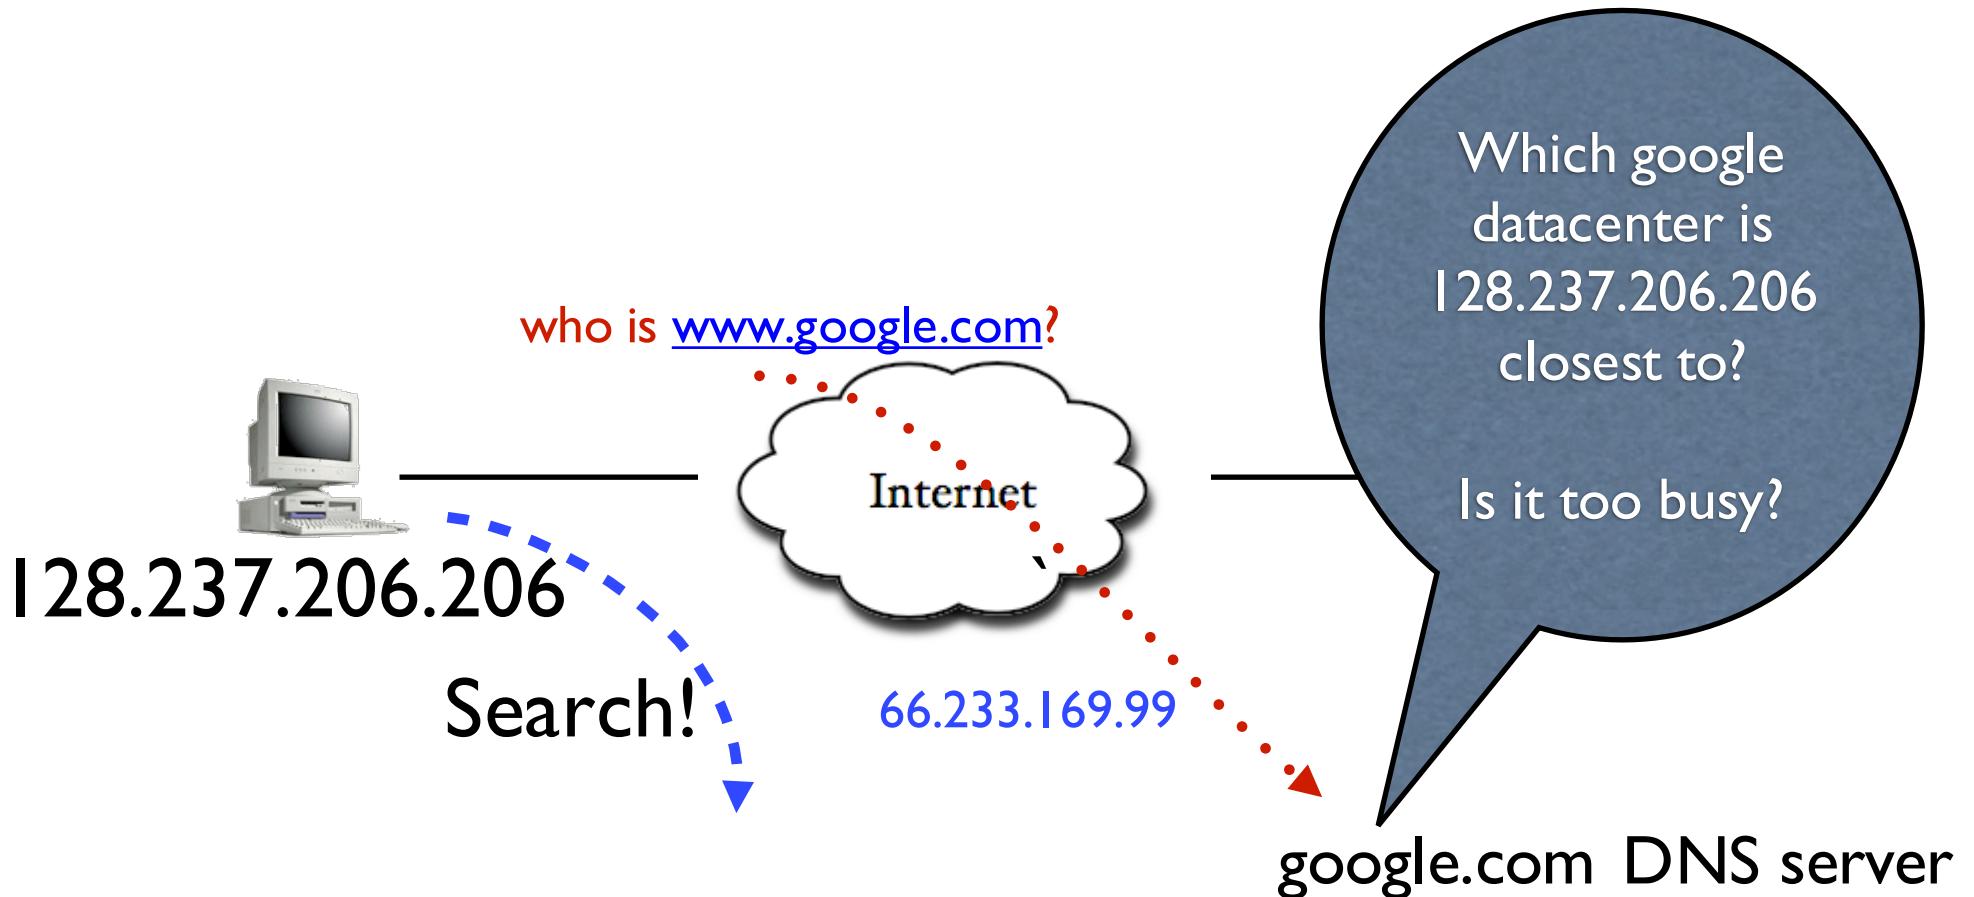
The Google Example

- Note that URL: www.google.com
- But your computer has an IP address...
- Naming! The “Domain Name System”, or DNS, translates names to IP addresses
 - In the days of yore, this was a text file called “hosts.txt” that everyone periodically downloaded
 - Today, with hundreds of millions of domains...
 - It’s a big distributed system that allows people to update small parts (“moo.cmcl.cs.cmu.edu”) without coordinating with the owners of other parts. We’ll see this soon.

Domain Name System



But there's more...



A Google Datacenter





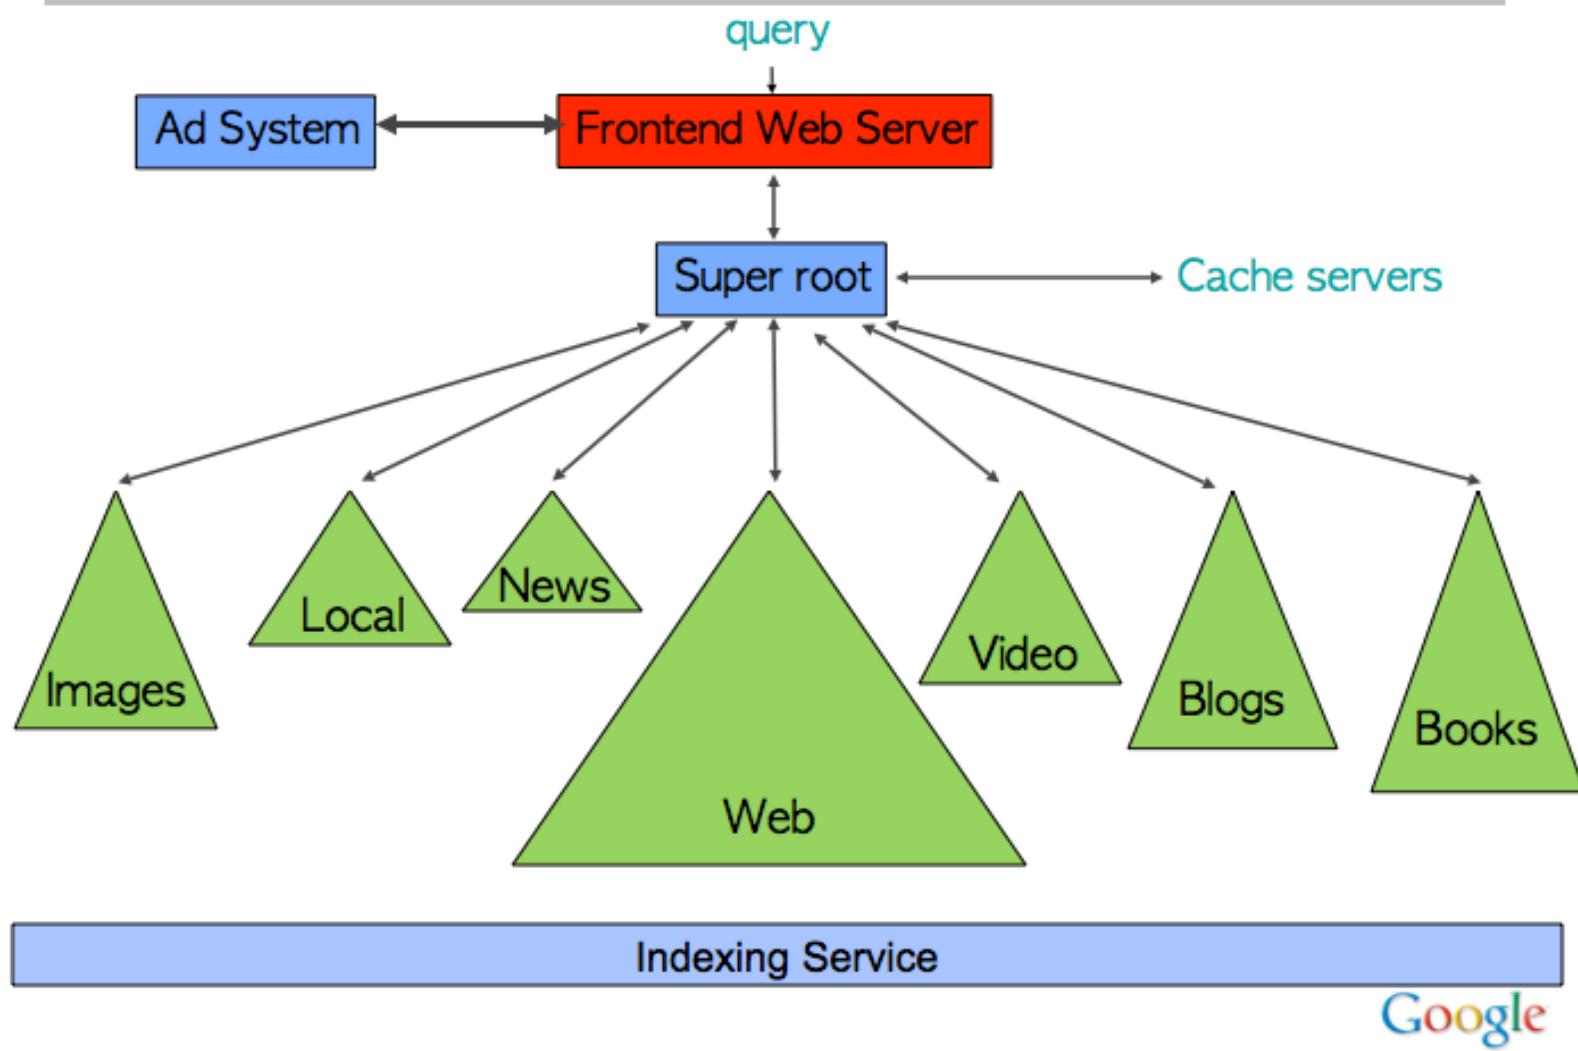
How big? Perhaps one million+ machines
but it's not that bad...
usually don't use more than 20,000 machines to
accomplish a single task. [2009, probably out of
date]

Search for Denali

Front-end



2007: Universal Search

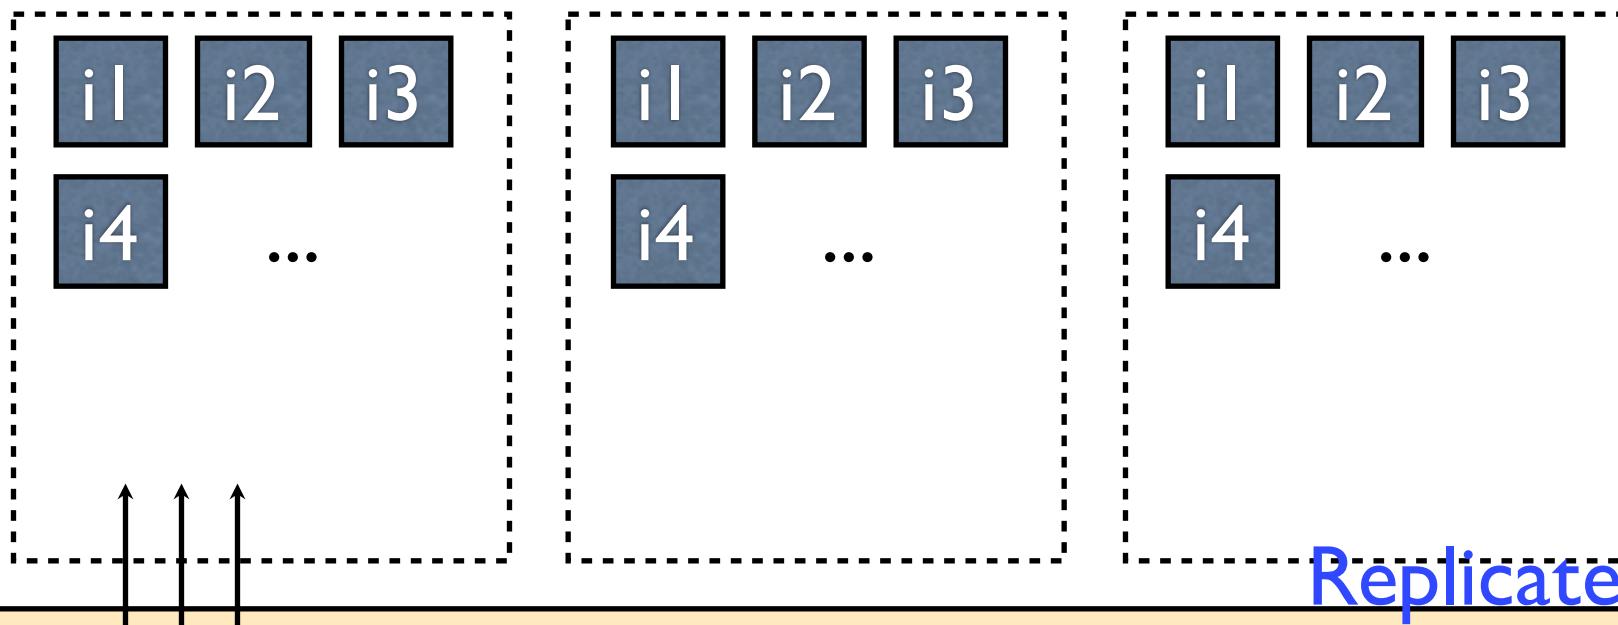


slide from Jeff Dean, Google

Front-end

Split into chunks:
make single
queries faster

Replicate:
Handle load



GFS distributed filesystem

Replicated
Consistent
Fast

How do you index the web?

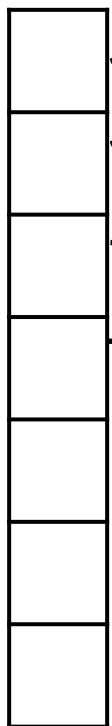
1. Get a copy of the web.
 2. Build an index.
 3. Prune.
- There are over 1 trillion unique URLs
- Billions of unique web pages
- Hundreds of millions of websites
- 30?? terabytes of text

==

- *Crawling* -- download those web pages
- *Indexing* -- harness 10s of thousands of machines to do it
- *Profiting* -- we leave that to you.
- “*Data-Intensive Computing*”

MapReduce / Hadoop

Data Why? Hiding details of programming 10,000
Chunl machines!



-Programmer writes two simple functions:

map (data item) -> list(tmp values)

reduce (list(tmp values)) -> list(out values)

...
MapReduce system balances load, handles
failures, starts job, collects results, etc.

Storage

Transformation Aggregation

All that...

- Hundreds of DNS servers
- Protocols on protocols on protocols
- Distributed network of Internet routers to get packets around the globe
- Hundreds of thousands of servers
- ... to find why Obama renamed a mountain?

Thanks!

Tuesday September 08, 1992

